

# Isolated and Distributed BGP attacks, and RPKI – From a RouteViews perspective

Kevin Conte  
University of Oregon  
Eugene, OR  
kconte@uoregon.edu

## ABSTRACT

abstract

## 1 INTRODUCTION

Many of the protocols that make the Internet work are chock-full of security flaws, and the Border Gateway Protocol is no different. It was developed in an era of the Internet absent of security concerns and knowledge. As a result, it is susceptible to many different kinds of attacks. To list a few, BGP hijacking, misconfigurations, black holes, and denial of service attacks (DoS).

A BGP hijacking is when a malicious actor advertises ownership of an Internet Protocol (IP) prefix that it does not actually own[1]. As a result, traffic is routed to and/or through this actor when it normally would not have.

Misconfigurations are similar, except they are an accident [11]. People make mistakes all the time. Often, these misconfigurations are due to "fat-finger" errors, where the maintainer of an Autonomous System (AS) mistypes important information, leading to incorrect and invalid advertisements.

Black holes are interesting in that they can be used both maliciously and assistingly. As far as malicious: black holes are a specific kind of hijack wherein all traffic that is routed to a black hole is simply lost. All packets entering the black hole are dropped and never reach their destination. Imagine a user is attempting to access a website whose server is under a prefix that is being routed to a black hole. From the perspective of a user, the site will seem to never load. This is because all traffic is lost. In the case of assists, BGP black holes can be used to mitigate DoS attacks. If one's organization is under attack, a BGP black hole can be used to direct all suspicious traffic away from the target[9, 10].

BGP Denial of Service attacks are quite similar to normal DoS attacks, just with BGP instead of a different target. These are identified through increased, abnormal traffic to a particular Autonomous System. This can be achieved by altering a routing table to direct traffic towards an AS.

Many of these attacks are identified through Invalid Route Origin advertisements. Unfortunately, there is no way to distinguish between such attacks without the use of extensive heuristics and context. The current "state of the art"

solution to this is developed and maintained by the Center for Applied Internet Data Analysis (CAIDA) and is called BGPStream [7]. Looking at historical data, it is incredibly difficult to distinguish between events as the historical context surrounding an event is not necessarily saved.

But, how does one defend against such attacks? Enter, Resource Public Key Infrastructure (RPKI). More detail is provided in Section 2.2. Simplifying, RPKI is the solution to BGP security[3], or lack thereof. Relying on trust anchors, RPKI provides route prefix validations to aid in the construction of accurate and safe routing tables.

In this study, I attempt to draw the line between isolated and distributed BGP attacks, and their respective trends over-time. To do this, I use the RouteViews dataset for BGP historical data, and RIPE NCC's dataset for historical RPKI data.

## 2 BACKGROUND

In this section I provide background for BGP and RPKI.

### 2.1 BGP

The Border Gateway Protocol (BGP) is the protocol that is used to communicate between discrete Autonomous Systems (ASes) [2]. It was developed to make the Internet more robust against incorporating individual networks into the whole Inter-network. BGP consists of announcements between peers. Peers are simply machines that communicate with one another; providing routing information about Internet Protocol (IP) prefixes. Peers will advertise all of their known prefixes to the other peers. These advertisements contain valuable information that influence the routing tables of the other peers in other Autonomous Systems.

The information contained in a BGP advertisement is as follows: timestamp, peer AS number (ASN), peer IP address, prefix, prefix length (in bits), AS Path, Next Hop, Origin AS.

The ASN of the peer is merely the AS that the peer belongs to. It is not necessarily the same ASN that owns the advertised prefix. The prefix is one of the most important parts of a BGP advertisement. It is what the peer is advertising it knows how to get to. The prefix length is the number of bits required to make subnet mask of that prefix. For example, the prefix 127.45.0.0/16 is a prefix of length 16 bits, so any

address of 127.45.0.0  $\rightarrow$  127.45.255.255 will match to that prefix. The AS Path is the shortest number of ASes that a packet must travel through to reach a destination prefix. The Next Hop attribute is the next IP address that a packet must be routed to in order to reach a particular prefix. It is often the same as the advertising peer's IP address, but not necessarily. Finally, the Origin AS is the AS that is advertising ownership of a prefix. This is the source of BGP violations.

Simplifying greatly, a BGP advertisement looks something like:

Peer ASN, Peer IP, Prefix, AS\_PATH, Origin AS  
1234, 192.56.23.10, 1.22.8.0/23, 1234 .. 45528, 45528

In this example, the peer is advertising from AS 1234. It is advertising the prefix: 1.22.8.0/23. And the originating AS is 45528.

For the purposes of this paper, focus on the Prefix and the Origin AS. This will become clearer in Section 4, when the methodology is described.

## 2.2 RPKI

The Resource Public Key Infrastructure (RPKI) [3] was developed by the Internet Engineering Task Force (IETF) in 2008 and saw initial release in 2011. It was developed to combat the security vulnerabilities of BGP. However, it has seen slow adoption by the Internet community, with Route Origin Validation not occurring until 2015.

RPKI provides objects called Route Origin Authorizations (ROAs). These objects provide validation of ownership for a particular prefix. Such objects contain the following information: ASN, prefix, maximum length, not before, not after. The ASN and prefix indicate that a particular AS owns that particular prefix. Maximum Length is used to aggregate multiple ROAs into a single ROA, for performance reasons. For example, if you are an organization that advertises two /24 prefixes under the same /23 prefix, it might be more efficient to have a single ROA that states your ownership of a /23 prefix, with the Maximum Length attribute set to 24. This can be a security issue when, if two different organizations own a /24 under a /23 but one of them has a ROA advertising ownership of the /23, all traffic gets routed to the "owner"; leaving the other without any valid traffic to their prefix. Finally, the not before and not after attributes of a ROA indicate the timespan that an AS owns a prefix. They may or may not own it before or after the timespan, but the indicated timespan guarantees ownership for that time.

ROAs, when validated are called Validated ROA Payloads (VRPs). All of the data used in this study come from VRPs, so I will henceforth refer to ROAs and VRPs interchangeably. Note that in a different context, they are different. ROAs

## RPKI hierarchy




Figure 1: RPKI Hierarchy [4]

are validated by what is known as a Trust Anchor Location (TAL). Since much of RPKI is dependent on *valid* and *trusted* ROAs, there are only a handful of TALs: IANA (Internet Assigned Numbers Authority), AFRINIC (African Network Information Centre), APNIC (Asia-Pacific Network Information Centre), ARIN (American Registry for Internet Numbers), LACNIC (Latin America and Caribbean Network Information Centre), and RIPE NCC (Réseaux IP Européens Network Coordination Centre).

## 3 DATASETS

In this section I provide an overview of the datasets used in this paper, and address the biases.

### 3.1 RouteViews

Historical BGP data comes courtesy of the University of Oregon's project: RouteViews. RouteViews contains full snapshots of what a particular collector has collected every two hours; and updates (or changes since last dump) every fifteen minutes. Historical data goes back to the 1990s, but only data from 2011 onward was used in this study. All RouteViews data is stored in file of the Multi-threaded Routing Toolkit file format[5]. Figure 2 shows the distribution of RouteViews collectors across the world. Each dot is a collector's location. RouteViews currently contains thirty-one collectors. It should be noted that this dataset is biased. When taking measurements on the internet, vantage points matter. RouteViews has a relatively low number of collectors worldwide, and a limited perspective of the Internet as a result. Thus, it cannot be said that this study is representative of the entire Internet. It can be said, however, to represent RouteViews's perspective of the Internet.

### 3.2 ROAs

Historical ROA data comes courtesy of RIPE NCC. They maintain a daily archive of RPKI dating back to its first deployment in 2011. It should be noted that there are several

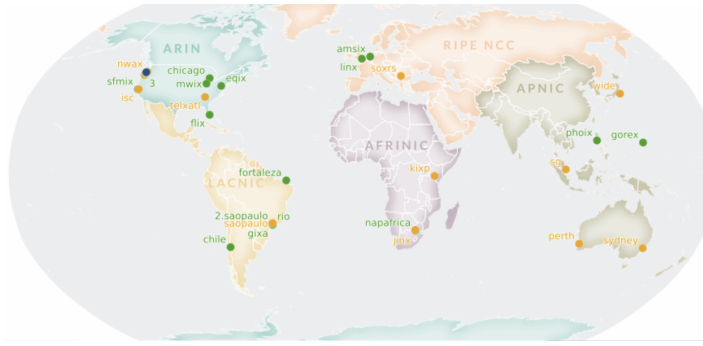


Figure 2: RouteViews Visibility Map

days which RIPE NCC does not have data for, due to a variety of reasons. Those days have, as such, been removed from this study.

## 4 METHODOLOGY

In this section I describe the methodology used.

The first step in deciphering the problem is to look at RPKI growth over time. Figure 3 shows the growth of RPKI in terms of the number of VRPs, beginning at RPKI's conception in 2011.

The second step is to gather BGP data and analyze it. I've taken over 1,600 samples between 21 January 2011 and 29 February 2020. Samples were taken once a day, every other day.

Next, to distinguish between isolated and distributed attacks I did the following. Define an **isolated** attack as an attack wherein *exactly* two discrete ASes advertise ownership of a particular prefix. Define a **distributed** attack as an attack wherein *more than* two discrete ASes advertise ownership of a particular prefix. Such attacks (isolated and distributed) will be called invalid route origins.

### 4.1 Assumptions

The following assumptions should be noted:

- (1) All "attacks" are just that, attacks. This study does not consider misconfigurations in its results. All invalid route origins are considered to be an attack of some sort. This is because, as mentioned earlier, invalid route origin causes are incredibly difficult to distinguish with any kind of accuracy.
- (2) All advertisements with more than one Origin ASN are considered to contain the valid Origin ASN. It is possible that some attacks include one of:
  - no valid Origin ASN
  - more than one valid Origin ASN

For the purposes of this study, we are assuming that each prefix has one and only one valid Origin ASN and

Internet Protocol	Advertisements	Prefixes	Attacks
IPv4	14,508,423,286	2,866,018	196,427
IPv6	854,637,095	326,627	11,414

Table 1: Summary of Results




Figure 3: Deployment Growth of RPKI

every prefix advertised contains the valid Origin ASN in one of its advertisements.

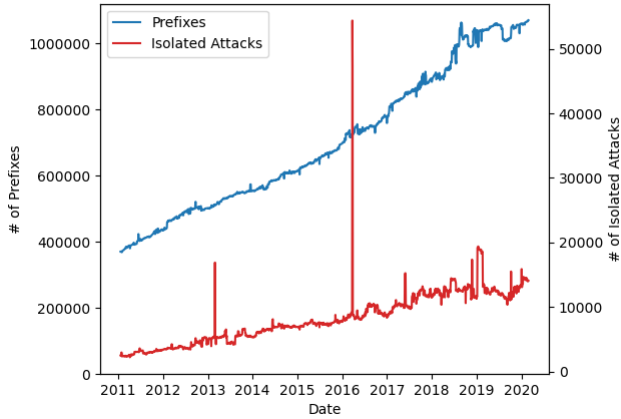
## 5 RESULTS

In this section I provide and discuss the results of the methodology.


### 5.1 RPKI Analysis

Again looking at Figure 3, the first thing to note is how long it took for RPKI to start validating. While such a service is critical to the security of the Internet, adoption clearly started out slow; with validation not occurring until 2015. However, as of August 2019, RPKI surpassed 100,000 VRPs. This is encouraging to the future success of RPKI. This is especially encouraging when one takes into account the apparent exponential adoption of RPKI.

Secondly, one may notice the spike of VRPs around June 2017. This was the result of one of the Regional Internet Registries, APNIC, transitioning to a new route management system. The transition caused the deaggregation of many ROAs from a few ASes [14]. Basically, there were a few ROAs with Max Length set to 24, while the prefix was much bigger than that. The ROAs were all valid, just re-configured.



**Figure 4: Number of Prefixes vs. Number of Isolated Attacks**



**Figure 5: Number of Prefixes vs. Percentage of Isolated Attacks**


## 5.2 Isolated Attacks

Almost 98% of the attacks seen in the RouteViews data were isolated attacks. Figure 4 shows the number of isolated attacks compared to the number of prefixes advertised over time. While both are growing, RPKI is having an effect. The gap between the number of isolated attacks and the number of prefixes advertised is growing. Figure 5 shows the percentage of isolated attacks over time compared to the number of prefixes advertised. Again, the gap between the two lines is growing. I'll also note that the spikes of the number of isolated attacks correlates very strongly to the spikes in the percentage of isolated attacks. What this means is that during the spikes, there is not an increase in the number of prefixes advertised. Rather, this is likely due to a coordinated attack against many prefixes. This is addressed further in Section 7.

## 5.3 Distributed Attacks

The remaining attacks are clearly distributed attacks. The biggest difference I noticed between distributed and isolated attacks is the volatility of distributed attacks. The counts and percentages vary wildly between timestamps. This indicates (coupled with the percentages) that isolated attacks are: a) much more common and b) likely easier to pull off. Again, the percentages of distributed attacks and the number of distributed attacks are strongly correlated. This draws the same conclusion as Section 5.2. Figures 6 and 7 provide similar information as Figures 4 and 5, just with distributed attacks.

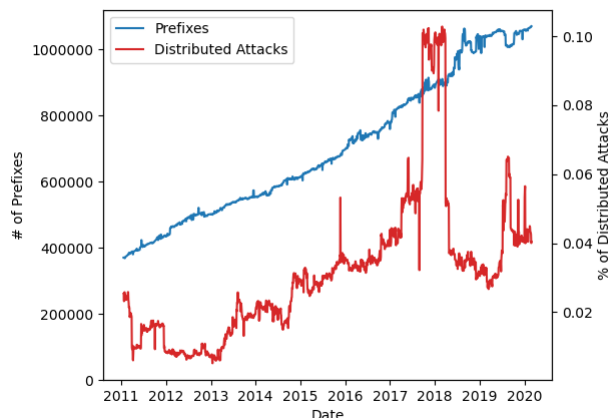
Figure 8 shows a CDF of the number of actors in a distributed attack. Most of the distributed attacks are occur with 5 or less actors, with around 80% occurring with 3 actors.




**Figure 6: Number of Prefixes vs. Number of Distributed Attacks**

## 6 CONCLUSIONS

From my findings, it is apparent that RPKI is having an effect, just not a huge one. The percentage of isolated attacks remains fairly constant throughout the study. It should be noted that my work is **not** peer reviewed and likely contains many errors both in methodology and results. The (lack of) effectiveness of RPKI is likely due to the vantage points from which my samples were taken. As stated earlier, RouteViews, while extensive, is not comprehensive. The vantage points I measured from could simply just not have been under the purview of RPKI coverage. This is entirely possible, considering that RouteViews only has 31 collectors worldwide, and



**Figure 7: Number of Prefixes vs. Percentage of Distributed Attacks**



**Figure 8: CDF of Number of Actors in an Attack**

even less during the duration of the study. As Table 1 shows, there were only about three million unique prefixes seen during the study. Considering that most advertisements are /24 IPv4 prefixes, three million is a small fraction of the total IPv4 /24 space (about sixteen million). And, not every prefix in the three million is a /24 prefix. All of this is to say that RouteViews doesn't actually see much of the Internet. Hopefully this would change in the future.

## 7 FUTURE WORK

In this section I propose future work surrounding this topic.

If one reads [8], they will see that my findings apparently contradict those found in that paper. This could be the result of several cases.

- (1) I have made a crucial error somewhere in my study (moderately likely).
- (2) The researchers in [8] have made an error (less likely).
- (3) I don't have enough data to extrapolate as much as I have.
- (4) The RouteViews dataset is not comprehensive enough.

In any case, future work could include validating my findings.

If I were to continue on with this study I would like to do a few things. First, take more samples. Ideally, I'd like to sift through all of the RouteViews dataset. Realistically, I know that is not possible. But, perhaps I could start with taking daily samples rather than bi-daily. This would double the amount of data I am looking at, hopefully providing further insight to the problem. Second, if I were to incorporate other sources of data, such as those from RIPE RIS, I could get more insight into RPKI's effectiveness. I recognize that RouteViews is limited in its scope. So, incorporating more data could help generate a representative study of the entire Internet and not just from the viewpoints measured.

I'd also like to perform some case studies. For example, I'd like to analyze the spikes on the various attack graphs to see what was happening at that time (if such information is available).

## 8 RESOURCES

All code used can be found at <https://github.com/Babarm/cis410-bgp-rpki-project>. This repository also contains the slides from my presentation on 9 March 2020.

As far as other tools used, I used BGPStream [6] (specifically the standalone binary, bgpreader) for parsing historical BGP data. For RPKI statistics, I used two tools, Routinator [12] for validating ROAs and Ziggy [13] for facilitating the validation of historical ROAs.

## 9 ACKNOWLEDGEMENTS

In this section I thank several groups and people for their help and contributions to this project. First, the authors of [8] for laying the ground work upon which this paper was possible. RouteViews for maintaining their various collectors and for providing their data for public use. RIPE NCC for maintaining the daily rsync repositories of ROA data across all five RIRs. CAIDA for developing and maintaining BGPStream [6] and for providing this tool for public use. NLnet Labs for developing and maintaining Ziggy [13] and Routinator [12] and for providing these tools for public use. I'd also like to thank Professor Ramakrishnan Durairajan at the University of Oregon for guiding me through this process. As well as David Teach, Senior Network Engineer at

the University of Oregon, and Steve Huter, Director of the Network Startup Resource Center, for allowing me to pick their brains while I was working on this project.

## REFERENCES

- [1] 2020. BGP hijacking. (2020). [https://en.wikipedia.org/wiki/BGP\\_hijacking](https://en.wikipedia.org/wiki/BGP_hijacking)
- [2] 2020. Border Gateway Protocol. (2020). [https://en.wikipedia.org/wiki/Border\\_Gateway\\_Protocol](https://en.wikipedia.org/wiki/Border_Gateway_Protocol)
- [3] 2020. Resource Public Key Infrastructure. (2020). [https://en.wikipedia.org/wiki/Resource\\_Public\\_Key\\_Infrastructure](https://en.wikipedia.org/wiki/Resource_Public_Key_Infrastructure)
- [4] 2020. RPKI Hierarchy. (2020). [https://blog.apnic.net/wp-content/uploads/2019/06/rpki\\_hierarchy.png](https://blog.apnic.net/wp-content/uploads/2019/06/rpki_hierarchy.png)
- [5] L. Blunk et al. 2011. *Multi-Threaded Routing Toolkit (MRT) Routing Information Export Format*. RFC 6396. RFC Editor. <https://www.rfc-editor.org/rfc/rfc6396.html>
- [6] CAIDA. 2020. BGPReader. (2020). <https://bgpstream.caida.org/docs/tools/bgpreader>
- [7] CAIDA. 2020. BGPStream. (2020). <https://bgpstream.com>
- [8] Taejoong Chung et al. 2019. RPKI is Coming of Age: A Longitudinal Study of RPKI Deployment and Invalid Route Origins. In *Proceedings of the Internet Measurement Conference*. 406–419.
- [9] David Davis. 2008. Black hole routes: The good, the bad, and the ugly. (2008). <https://www.techrepublic.com/blog/data-center/black-hole-routes-the-good-the-bad-and-the-ugly/>
- [10] Mattijs Jonker et al. 2018. A First Joint Look at DoS Attacks and BGP Blackholing in the Wild. In *Proceedings of the Internet Measurement Conference*.
- [11] Ratul Mahajan et al. 2002. Understanding BGP Misconfiguration.
- [12] NLnetLabs. 2019. Routinator. (2019). <https://nlnetlabs.nl/projects/rpki/routinator>
- [13] NLnetLabs. 2019. Ziggy. (2019). <https://github.com/NLnetLabs/ziggy>
- [14] Roland van Rijswijk-Deij. 2020. Leaping through RPKI history with Ziggy. (August 2020). <https://blog.apnic.net/2019/08/30/leaping-through-rpki-history-with-ziggy/>